

SC15-009: Recent Advances in Physics-Informed Deep Learning

Physics-Informed Deep Generative Models


Paris Perdikaris

Department of Mechanical Engineering
University of Pennsylvania
email: pgp@seas.upenn.edu

USNCCM15
Austin, TX
July 28, 2019



Data-driven modeling of stochastic systems



Physics-informed deep generative models: $p(u|x, t, z), z \sim p(z)$, such that $u_t + \mathcal{N}_x u = 0$.

Current approaches (non-intrusive methods):

- Polynomial chaos, sparse grid quadratures, multi-level/multi-fidelity Monte Carlo, reduced order/surrogate models (POD, Gaussian processes, etc.)
- All face limitations in modeling high-dimensional stochastic systems.
- All require repeated evaluations of expensive simulators/experiments.


Goal of this work:

- Introduce a probabilistic deep learning framework for modeling stochastic systems that entirely bypasses the need for repeatedly sampling expensive experiments or numerical simulators.

Approach:

- Build deep generative models (GANs, VAEs, etc.) with physics-informed constraints.
- Develop robust statistical inference algorithms for approximating complex conditional distributions directly from noisy data.


Conditional deep generative models



Model: $\mathbf{u} = f_\theta(\mathbf{x}, t, \mathbf{z}), \quad \mathbf{z} \sim p(\mathbf{z}),$ such that $\mathbf{u}_t + \mathcal{N}_{\mathbf{x}}\mathbf{u} = 0$

$$p_\theta(\mathbf{u}|\mathbf{x}, t) = \int p_\theta(\mathbf{u}, \mathbf{z}|\mathbf{x}, t) d\mathbf{z} = \int p_\theta(\mathbf{u}|\mathbf{x}, t, \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$

Training: $\text{KL}[p_\theta(\mathbf{x}, t, \mathbf{u}) || q(\mathbf{x}, t, \mathbf{u})]$



Key ingredients:

1. Density ratio estimation via probabilistic classification.
2. Joint distribution matching via adversarial inference with entropy regularization.
3. Physics-informed constraints for generating samples that approximately satisfy the underlying PDE.


#1: Density ratio estimation via probabilistic classification

$$\mathbb{KL}[p_\theta(\mathbf{x}, t, \mathbf{u}) || q(\mathbf{x}, t, \mathbf{u})] := \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u})} \left[\log \left(\frac{p_\theta(\mathbf{x}, t, \mathbf{u})}{q(\mathbf{x}, t, \mathbf{u})} \right) \right]$$

Estimating density ratios is a challenging task:


- Each part of the ratio may itself involve intractable integrals
- We often deal with high-dimensional quantities.
- We may only have samples drawn from the two distributions, not their analytical forms.

This is where the **density ratio trick** enters:
it allows us to construct a binary classifier
that distinguishes between samples from the
two distributions.



The density ratio gives the correction factor needed to make two distributions equal.

$$\begin{aligned}
 \frac{p_\theta(\mathbf{x}, t, \mathbf{u})}{q(\mathbf{x}, t, \mathbf{u})} &= \frac{\rho(\mathbf{x}, t, \mathbf{u}|y = +1)}{\rho(\mathbf{x}, t, \mathbf{u}|y = -1)} \\
 &= \frac{\rho(y = +1|\mathbf{x}, t, \mathbf{u})\rho(\mathbf{x}, t, \mathbf{u})}{\rho(y = +1)} \Big/ \frac{\rho(y = -1|\mathbf{x}, t, \mathbf{u})\rho(\mathbf{x}, t, \mathbf{u})}{\rho(y = -1)} \\
 &= \frac{\rho(y = +1|\mathbf{x}, t, \mathbf{u})}{\rho(y = -1|\mathbf{x}, t, \mathbf{u})} = \frac{\rho(y = +1|\mathbf{x}, t, \mathbf{u})}{1 - \rho(y = +1|\mathbf{x}, t, \mathbf{u})} \\
 &= \frac{T(\mathbf{x}, t, \mathbf{u})}{1 - T(\mathbf{x}, t, \mathbf{u})}.
 \end{aligned}$$




#2: Joint distribution matching

- Physics-informed generative model : $p(\mathbf{u}|\mathbf{x}, t, \mathbf{z})$, $\mathbf{z} \sim p(\mathbf{z})$, such that $\mathbf{u}_t + \mathcal{N}_{\mathbf{x}}\mathbf{u} = 0$.
- We train the model via joint distribution matching by minimizing the reverse KL-divergence :

$$\text{KL}[p_\theta(\mathbf{x}, t, \mathbf{u})||q(\mathbf{x}, t, \mathbf{u})] = -h(p_\theta(\mathbf{x}, t, \mathbf{u})) - \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u})}[\log(q(\mathbf{x}, t, \mathbf{u}))]$$

$$= \underbrace{-h(p_\theta(\mathbf{x}, t, \mathbf{u})))}_{\text{spreads the support of } p_\theta} - \int_{\mathcal{S}_{p_\theta} \cap \mathcal{S}_q} \log(q(\mathbf{x}, t, \mathbf{u})) p_\theta(\mathbf{x}, t, \mathbf{u}) dx dt d\mathbf{u} - \underbrace{\int_{\mathcal{S}_{p_\theta} \cap \mathcal{S}_q^o} \log(q(\mathbf{x}, t, \mathbf{u})) p_\theta(\mathbf{x}, t, \mathbf{u}) dx dt d\mathbf{u}}_{\text{penalizes non-overlaps of } p_\theta \text{ and } q}$$



$p_\theta(\mathbf{x}, t, \mathbf{u})$: Generative model distribution

$q(\mathbf{x}, t, \mathbf{u})$: Empirical data distribution

- Variational bound for the intractable entropy term :
$$\begin{aligned} h(p_\theta(\mathbf{x}, t, \mathbf{u})) &= h(p(\mathbf{z})) - h(p_\theta(\mathbf{z}|\mathbf{x}, t, \mathbf{u})) + \cancel{h(p_\theta(\mathbf{x}, t, \mathbf{u}|\mathbf{z}))} \rightarrow 0 \\ &= h(p(\mathbf{z})) + \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u}, \mathbf{z})}[\log(p_\theta(\mathbf{z}|\mathbf{x}, t, \mathbf{u}))] \\ &= h(p(\mathbf{z})) + \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u}, \mathbf{z})}[\log(q_\phi(\mathbf{z}|\mathbf{x}, t, \mathbf{u}))] \\ &\quad + \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u})}[\text{KL}[p_\theta(\mathbf{z}|\mathbf{x}, t, \mathbf{u})||q_\phi(\mathbf{z}|\mathbf{x}, t, \mathbf{u})]] \\ &\geq h(p(\mathbf{z})) + \mathbb{E}_{p_\theta(\mathbf{x}, t, \mathbf{u}, \mathbf{z})}[\log(q_\phi(\mathbf{z}|\mathbf{x}, t, \mathbf{u}))]. \end{aligned}$$

Li, C., Li, J., Wang, G., & Carin, L. (2018). Learning to Sample with Adversarially Learned Likelihood-Ratio.

Remarks:


- Flexible variational inference with implicit distributions, no mean field approximations.
- Explicit control over the pathology of mode-collapse in GANs.
- Approximate posterior inference over latent variables (not possible with GANs).
- Discovery of disentangled representations via cycle-consistency in latent space.

#3: Physics-informed constraints

$f_\theta(\mathbf{x}, t)$: Neural network

$r_\theta(\mathbf{x}, t)$: Physics-informed neural network

$$\begin{cases} [\mathbf{x}, t] \xrightarrow{f_\theta} \mathbf{u}(\mathbf{x}, t) \\ [\mathbf{x}, t] \xrightarrow{r_\theta} \frac{\partial}{\partial t} f_\theta(\mathbf{x}, t) + \mathcal{N}_\mathbf{x} f_\theta(\mathbf{x}, t) \end{cases}$$



Lagaris, I. E., Likas, A., & Fotiadis, D. I. (1998). Artificial neural networks for solving ordinary and partial differential equations. *IEEE transactions on neural networks*, 9(5), 987-1000.

Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686-707.

$$\mathcal{L}_{\text{PDE}}(\theta) := \underbrace{\frac{1}{N_r} \sum_{I=1}^{N_r} \|r_\theta(\mathbf{x}, t) - \mathbf{r}_i\|^2}_{\text{physics-informed regularization}}$$

Adversarial inference for physics-informed deep generative models

- Adversarial optimization:

$$\max_{\psi} \mathcal{L}_{\mathcal{D}}(\psi)$$

$$\min_{\theta, \phi} \mathcal{L}_{\mathcal{G}}(\theta, \phi) + \beta \mathcal{L}_{PDE}(\theta),$$

$p(z) \rightarrow \text{prior}$
 $f_{\theta}(x, t, z) \rightarrow \text{generator}$
 $T_{\psi}(x, t, u) \rightarrow \text{discriminator}$
 $q_{\phi}(z|x, t, u) \rightarrow \text{encoder}$

- Generator and discriminator loss functions:

$$\mathcal{L}_{\mathcal{D}}(\psi) = \mathbb{E}_{q(x, t)p(z)}[\log \sigma(T_{\psi}(x, t, f_{\theta}(x, t, z)))] + \mathbb{E}_{q(x, t, u)p(z)}[\log(1 - \sigma(T_{\psi}(x, t, f_{\theta}(x, t, z))))]$$

$$\mathcal{L}_{\mathcal{G}}(\theta, \phi) = \mathbb{E}_{q(x, t, u)p(z)}[T_{\psi}(x, t, f_{\theta}(x, t, z)) + (1 - \lambda) \log(q_{\phi}(z|x, t, f_{\theta}(x, t, z)))]$$

- PDE constraints:

$$\mathcal{L}_{PDE}(\theta) := \frac{1}{N_r} \sum_{I=1}^{N_r} \|r_{\theta}(x, t) - r_i\|^2$$


physics-informed regularization

- Alternate stochastic gradient updates between the generator and the discriminator.

- Optimal discriminator loss:

$$\ln(4) = -2 \times \ln(0.5) = 1.384.$$

Stochastic gradient descent dynamics



A pedagogical example

Problem setup:

$$u_{xx} - u^2 u_x = f(x), \quad x \in [-1, 1],$$

$$f(x) = -\pi^2 \sin(\pi x) - \pi \cos(\pi x) \sin^2(\pi x)$$


$$u(-1), u(1) \sim \mathcal{N}(\mathbf{0}, \sigma_n^2 \mathbf{I})$$

Model setup and training data:


- 20 realizations at each boundary point
- 100 collocation points for enforcing the PDE residual
- $\lambda = 1.5, \beta = 1.0$

Neural nets: Feed-forward with 2 hidden layers, 50 neurons, $\tanh()$ activation, Adam optimizer.

$\sigma_n^2 = 0.0$ (deterministic case)



$\sigma_n^2 = 0.05$ (stochastic case)



Mean and two standard deviations of $p_\theta(u|x, z)$


A pedagogical example

Sensitivity wrt λ, β


$$\mathbb{E}_{p(x)}\{\text{KL}[p_\theta(u|x)||q(u|x)]\}$$

$\beta \backslash \lambda$	1.0	1.5	2.0	5.0
0	5.0e+05	7.5e+01	6.0e+01	4.4e+01
1.0	3.3e+02	1.8e-01	2.9e-01	2.0e-01
2.0	2.1e+02	1.7e-01	5.0e-02	1.2e-01
5.0	3.5e+01	1.8e-01	1.9e-01	1.1e-01

Mode-collapse for $\lambda = 1.0$



Stable prediction for $\lambda = 1.5$



Uncertainty propagation in nonlinear conservation laws


Problem setup:

Burgers equation

$$\begin{cases} u_t + uu_x - \nu u_{xx} = 0, \\ u(0, x) = -\sin(\pi x), \\ u(t, -1) = u(t, 1) = 0, \\ x \in [-1, 1], t \in [0, 1], \\ \nu = 0.01/\pi \end{cases}$$

Model setup and training data:


- 100 scattered measurements from a random initial condition
- 10,000 collocation points for enforcing the PDE residual
- $\lambda = 1.5, \beta = 1.0$



Random initial condition:


$$u(x, 0) = -\sin(\pi(x + 2\delta)) + \delta,$$

$$\delta = \frac{\epsilon}{\exp(3|x|)}, \quad \epsilon \sim N(0, 0.1^2)$$



Neural nets: Feed-forward with 4 hidden layers, 50 neurons, `tanh()` activation, Adam optimizer.

Summary



Physics-informed deep generative models: $p(\mathbf{u}|\mathbf{x}, t, z)$, $z \sim p(z)$, such that $\mathbf{u}_t + \mathcal{N}_{\mathbf{x}}\mathbf{u} = 0$.

Advantages:

- Approximate arbitrarily complex and high-dimensional probability distributions.
- Bypasses the need for repeatedly sampling expensive experiments or numerical simulators.
- Encourage generative models to produce samples that satisfy PDEs.
- Avoid over-simplifying approximations (e.g. mean-field variational inference).
- Enables general and flexible schemes for statistical inference

Caveats:

- Adversarial models require careful tuning.
- Theoretical asymptotic behavior is hard to be achieved in practice.
- Further understanding needs to be gained before tackling cases exhibiting multi-scale dynamics, chaos, turbulence, etc.

Acknowledgements:



Yibo Yang
(Penn)



U.S. DEPARTMENT OF
ENERGY



Hands-on tutorial:

<https://github.com/PredictiveIntelligenceLab/USNCCM15-Short-Course-Recent-Advances-in-Physics-Informed-Deep-Learning/blob/master/notebooks/PIDGMs.ipynb>

- Yang, Y., & Perdikaris, P. (2018). Physics-informed deep generative models. *Neural Information Processing Systems, Workshop on Bayesian Deep Learning*.
- Yang, Y., & Perdikaris, P. (2019). Adversarial uncertainty quantification in physics-informed neural networks. *Journal of Computational Physics*.

Code: <https://github.com/PredictiveIntelligenceLab/UQPINNs>